

# ECE 459: Programming for Performance

## Assignment 3

Zhidong Zhang

March 10, 2020

### Part 1: Crack it to me

#### 1. Design Choices:

- If the secret length is  $x$  and the number of possible characters is 36, then the number of possible secrets is  $36^x$ . For the secret length of  $x$ , the problem size is  $36^x$ , which can be set as the global work size of the kernel. In the kernel, `get_global_id(0)` returns a global id from 0 to  $36^x - 1$ . We can create a Base36 decoder to decode the global id into a possible secret with length  $x$ . In the same kernel, I verify the correctness of the possible secret by comparing the original signature and the current signature generated by the HMAC\_SHA256 function. If the secret is correct, I will write the secret into the output buffer.
- In the `main.cpp`, the JWT secret's max length is given as the `gMaxSecretLen`. Therefore, possible secret lengths are ranging from  $[1 : gMaxSecretLen]$ . To find the secret, for each possible secret length  $x \in [1 : gMaxSecretLen]$ , I run the kernel separately with the global size set to  $36^x$ .

#### 2. Results:

	<i>jwtcracker_sin</i>	<i>jwtcracker_omp</i>	<i>jwtcracker_ocl</i>
<i>MaxSecretLen = 4</i>	2.593	<b>0.185s</b>	1.424s
<i>MaxSecretLen = 5</i>	102.06s	7.444s	<b>3.284s</b>
<i>MaxSecretLen = 6</i>	4104s	306s	<b>75s</b>

#### 3. Speed up explanation

- Here are the speedups observed from the above table. When the `maxSecretLen` is 4, `jwtcracker_omp` is the fastest. This is because `jwtcracker_ocl` spawning  $36^4$  number of threads whose overhead outweigh its speedup when the `maxSecretLen` is small. When the `maxSecretLen` is 5, `jwtcracker_ocl` is the fastest, which has 31 times speedup compared with `jwtcracker_ocl` and 2 times speedup compared with `jwtcracker_omp`. When the `maxSecretLen` is 6, `jwtcracker_ocl` is the fastest, which has 55 times speedup compared with `jwtcracker_ocl` and 4 times speedup compared with `jwtcracker_omp`. This is expected. When the secret length is longer, OpenCL shows a significant advantage over OpenMP, because all the threads or work items are distributed and parallelized by the 2560 CUDA Cores in the GPU.

## Part 2: Coulomb's Law Problem

### 1. Design Choices:

- In the OpenMP code, functions computeForces(), computeApproxPositions(), computeBetterPositions() and isErrorAcceptable() is applied with openMp parallel-for directive, which enables parallelization. Therefore, in the OpenCL open, I need to transfer those functions into different kernels functions. I created three kernel functions which are computeForces(), computePositions(), isErrorAcceptable(), whose global work size is set to the total number of particles. In each kernel, get\_global\_id(0) returns the global id which is the particle id or particle index. Each kernel uses the particle id to compute the property just for that particle.
- In Simulation::run() function, I substituted all openMP function calls with kernels calls, and created buffers and arguments correctly.

### 2. Results:

	<i>protons.sin</i>	<i>protons.omp</i>	<i>protons.ocl</i>
<i>numProtos = 25, numElections = 25</i>	<b>0.025s</b>	3.570s	1.114s
<i>numProtos = 500, numElections = 500</i>	4.541	4.809s	<b>1.186s</b>
<i>numProtos = 1000, numElections = 1000</i>	16.347s	5.905s	<b>1.532s</b>
<i>numProtos = 2500, numElections = 2500</i>	109.609s	17.624s	<b>2.075s</b>
<i>numProtos = 5000, numElections = 5000</i>	441.628s	56.321s	<b>3.145s</b>
<i>numProtos = 7500, numElections = 7500</i>	997.164s	109.303s	<b>4.265s</b>
<i>numProtos = 10000, numElections = 10000</i>	1998.533s	202.066s	<b>5.216s</b>

### 3. Speed up explanation

- From the above results, when the number of particles is small, the *protons.sin* works the best because it has the lowest overhead. However, when the number of particles is bigger (over 500), the *protons.ocl* shows a significant advantage over *protons.omp* and *protons.sin* by achieving more speed up as the number of particles increases. Especially, when the total number of elections is 10000, *protons.ocl* has around 400 times speedup over the single-thread solution and 40 times speedup over the OpenMP solution. This is as expected. When the number of particles is small, the GPU is under-utilized. As the number of particles increases by  $x$  time, the *protons.sin*'s and *protons.omp*'s execution time increase by  $x^2$  time, whereas, the *protons.ocl*'s execution time only increase by  $x$  time, because the outer loop of function *computeForce()* and *computerPositions()* are fully parallelized on the GPU.